

```
In [1]: import csv
import sys
from datetime import datetime
import pandas as pd
from statsmodels.tsa.stattools import grangercausalitytests
import pandas as pd
import numpy as np
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Dropout
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
```

2023-05-20 03:48:42.266214: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2  
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
In [2]: company_name = 'Cisco'
close_price_as_predictive_column = False

start_date = datetime.strptime('03/02/2020', '%m/%d/%Y')
end_date = datetime.strptime('03/03/2023', '%m/%d/%Y')

data = []

with open('Historical Data/' + company_name + '.csv', 'r') as csv_file:
    csv_reader = csv.reader(csv_file)
    next(csv_reader)

    for row in csv_reader:
        date_str = row[0]
        date = datetime.strptime(date_str, '%m/%d/%Y')
        if start_date <= date <= end_date:
            data.append(row)

rows = list(data)

averages = []
for i in range(0, len(rows), 5):
    row_subset = rows[i:i+5]
    col_subset = [float(row[1][1:]) for row in row_subset]
    avg = sum(col_subset) / len(col_subset)
    averages.append(round(avg,2))

processed_data = []
processed_data.append(['Date', 'Close', 'Layoffs'])
for i in range(0, len(rows), 5):
    row_subset = rows[i:i+5]
    first_row = row_subset[0][0]
    mapped_avg = averages[int(i/5)]
    date = datetime.strptime(first_row, '%m/%d/%Y')
    processed_data.append([first_row, mapped_avg,0])
```

```

In [3]: companies = ['Adobe', 'Airbnb', 'Amazon', 'Apple', 'Atlassian', 'Cisco', 'Co

with open('layoffs.csv', 'r') as csv_file:
    csv_reader = csv.reader(csv_file)
    rows = list(csv_reader)

    """
    for row in csv_reader:
        if row[0] in companies:
            print(row)
    """

sorted_rows = sorted(rows, key=lambda x: x[0])

processed_rows = list(processed_data)

start = 1

for row in sorted_rows:
    if (row[0] == company_name):
        print(row)
        date = datetime.strptime(row[5], '%Y-%m-%d')
        for i in range(start, len(processed_rows)-1):
            second_date = datetime.strptime(processed_rows[i][0], '%m/%d/%Y')
            first_date = datetime.strptime(processed_rows[i+1][0], '%m/%d/%Y')
            if second_date == date:
                processed_data[i][2] = int(row[3])
                start = i
            else:
                if first_date == date:
                    processed_data[i+1][2] = int(row[3])
                    start = i
                else:
                    if first_date < date < second_date and row[3].isdigit():
                        processed_data[i+1][2] = int(row[3])
                        start = i

with open('Processed Data/' + company_name + '.csv', 'w', newline='') as csv
    csv_writer = csv.writer(csv_file)
    for row in processed_data:
        csv_writer.writerow(row)

['Cisco', 'SF Bay Area', 'Infrastructure', '4100', '0.05', '2022-11-16', 'Po
st-IPO', 'United States', '2.0']

```

```
In [4]: df = pd.read_csv('Processed Data/' + company_name + '.csv')

ts1 = df.iloc[:, 1]
ts2 = df.iloc[:, 2]

results = grangercausalitytests(df[['Layoffs', 'Close']], maxlag=4, verbose=0)

print(results)

for lag in range(1, 5):
    print(f'Lag {lag}:')
    print(f'F-test p-value: {results[lag][0]["params_ftest"][1]}')
    print(f'Chi-squared p-value: {results[lag][0]["ssr_chi2test"][1]}')
    print('\n')
```

```
{1: ({'ssr_ftest': (0.19506614726493515, 0.6593777627979982, 148.0, 1), 'ssr_chi2test': (0.19902019079057573, 0.6555128848116205, 1), 'lrtest': (0.19888915017054387, 0.655618989966779, 1), 'params_ftest': (0.19506614726478214, 0.6593777627981194, 148.0, 1.0)}, [<statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefal370f10>, <statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefal370370>, array([[0., 1., 0.]])]), 2: ({'ssr_ftest': (1.8492376661899852, 0.16104890830951132, 145.0, 2), 'ssr_chi2test': (3.8260089645310043, 0.14763614938565114, 2), 'lrtest': (3.7780286520187474, 0.15122078980032644, 2), 'params_ftest': (1.8492376661899415, 0.16104890830951787, 145.0, 2.0)}, [<statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefd243dca0>, <statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefd243d850>, array([[0., 0., 1., 0., 0.]], [0., 0., 0., 1., 0.]])]), 3: ({'ssr_ftest': (1.2553797315704955, 0.2921240011913806, 142.0, 3), 'ssr_chi2test': (3.95179394374656, 0.26671678319343484, 3), 'lrtest': (3.900297537190454, 0.27243340912215636, 3), 'params_ftest': (1.2553797315704842, 0.2921240011913826, 142.0, 3.0)}, [<statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefd243d7c0>, <statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefd243dc70>, array([[0., 0., 0., 1., 0., 0., 0.]], [0., 0., 0., 0., 1., 0., 0.]], [0., 0., 0., 0., 0., 1., 0.]])]), 4: ({'ssr_ftest': (1.1276045345161594, 0.34605749923378787, 139.0, 4), 'ssr_chi2test': (4.802459600241485, 0.30817338133477945, 4), 'lrtest': (4.726187555447723, 0.3165631724820432, 4), 'params_ftest': (1.1276045345161352, 0.34605749923379747, 139.0, 4.0)}, [<statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefal3e40d0>, <statsmodels.regression.linear_model.RegressionResultsWrapper object at 0x7fefal3e41f0>, array([[0., 0., 0., 0., 1., 0., 0., 0., 0.]], [0., 0., 0., 0., 0., 1., 0., 0., 0.]], [0., 0., 0., 0., 0., 0., 1., 0., 0.]], [0., 0., 0., 0., 0., 0., 0., 1., 0.]])])}]
```

Lag 1:

F-test p-value: 0.6593777627981194

Chi-squared p-value: 0.6555128848116205

Lag 2:

F-test p-value: 0.16104890830951787

Chi-squared p-value: 0.14763614938565114

Lag 3:

F-test p-value: 0.2921240011913826

Chi-squared p-value: 0.26671678319343484

Lag 4:

F-test p-value: 0.34605749923379747

Chi-squared p-value: 0.30817338133477945

```
In [5]: df = df.iloc[:::-1]
df['Pct_Change'] = round(df['Close'].pct_change(periods=1)*100,2)
if df['Pct_Change'].isnull().values.any():
    df['Pct_Change'] = df['Pct_Change'].fillna(0)

def direction(row):
    if(row['Pct_Change'] < 0):
        return -1
    else:
        return 1

df['Direction_Change'] = df.apply(direction, axis=1)

#print(df)
df.to_csv('Processed Data/' + company_name + '.csv', index=False, mode='w')
```

```
In [6]: X = df[['Layoffs','Close']]
y = df['Close']
#y = df['Pct_Change']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

xgb_model = xgb.XGBRegressor(objective='reg:squarederror', colsample_bytree
    max_depth = 5, alpha = 10, n_estimators = 100)

xgb_model.fit(X_train, y_train)

y_pred = xgb_model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("RMSE:", rmse)

r2 = r2_score(y_test, y_pred)
print('R2 score: {:.2f}'.format(r2))

RMSE: 0.21648598470263977
R2 score: 1.00
```

```
In [7]: df = pd.read_csv('Processed Data/' + company_name + '.csv')
print(df)
```

	Date	Close	Layoffs	Pct_Change	Direction_Change
0	03/04/2020	40.87	0	0.00	1
1	03/11/2020	38.87	0	-4.89	-1
2	03/18/2020	35.43	0	-8.85	-1
3	03/25/2020	36.84	0	3.98	1
4	04/01/2020	39.47	0	7.14	1
..	...	...	...	...	...
147	02/02/2023	48.66	0	2.06	1
148	02/09/2023	47.55	0	-2.28	-1
149	02/16/2023	48.45	0	1.89	1
150	02/24/2023	49.49	0	2.15	1
151	03/03/2023	48.66	0	-1.68	-1

[152 rows x 5 columns]

```
In [8]: inputcols =['Layoffs','Close']
outputcols = ['Direction_Change']
data = df[inputcols]
scaler = StandardScaler()
scaled = scaler.fit_transform(data)
nlags = 4
X = []
y = []
for i in range(nlags, len(scaled)):
    X.append(scaled[i - nlags:i])
    if(close_price_as_predictive_column):
        y.append(scaled[i,1]) # for close price as a predictive column

if(not close_price_as_predictive_column):
    y = np.array(df.loc[nlags:,outputcols]) #for direction change as a predictive column
X = np.array(X)
y = np.array(y)

sc_predict = StandardScaler()
sc_predict.fit_transform(np.array(data.iloc[:, 1]).reshape(-1, 1))
```

```
Out[8]: array([[ -1.21066426],
[ -1.53480748],
[ -2.09233381],
[ -1.86381284],
[ -1.43756451],
[ -1.26576861],
[ -1.08586912],
[ -1.06155838],
[ -0.90434892],
[ -1.14907705],
[ -0.85410672],
[ -0.57372283],
[ -0.40840979],
[ -0.22364816],
[ -0.30306325],
[ -0.40678908],
```

```
[ -0.4797213 ],  
[ -0.35006401 ],  
[ -0.35492616 ],  
[ -0.30954611 ],  
[ -0.22040673 ],  
[ -0.25930391 ],  
[ -0.12154304 ],  
[ -0.61910288 ],  
[ -1.00807475 ],  
[ -0.99835045 ],  
[ -1.15069776 ],  
[ -1.33383868 ],  
[ -1.3857016 ],  
[ -1.58667039 ],  
[ -1.55911822 ],  
[ -1.37759802 ],  
[ -1.40028804 ],  
[ -1.63529187 ],  
[ -2.00319443 ],  
[ -1.73253484 ],  
[ -1.22849214 ],  
[ -1.10207628 ],  
[ -0.84438242 ],  
[ -0.64179291 ],  
[ -0.62882718 ],  
[ -0.56723997 ],  
[ -0.59479214 ],  
[ -0.62720647 ],  
[ -0.47323844 ],  
[ -0.52672207 ],  
[ -0.49916989 ],  
[ -0.18961312 ],  
[ -0.14909522 ],  
[ -0.40516836 ],  
[ -0.45865199 ],  
[ -0.33547757 ],  
[ 0.10697792 ],  
[ 0.18801373 ],  
[ 0.49594979 ],  
[ 0.58995132 ],  
[ 0.5478127 ],  
[ 0.62560707 ],  
[ 0.51377766 ],  
[ 0.44408687 ],  
[ 0.70502216 ],  
[ 0.71474646 ],  
[ 0.75202293 ],  
[ 0.74716078 ],  
[ 0.95461244 ],  
[ 0.89626666 ],  
[ 0.72447076 ],  
[ 0.76660937 ],  
[ 0.81361014 ],
```



```
[ 0.85250733],  
[ 0.92706027],  
[ 1.11020119],  
[ 1.20096129],  
[ 1.23499632],  
[ 1.33710144],  
[ 1.74228046],  
[ 1.76497049],  
[ 1.6353132 ],  
[ 1.41975796],  
[ 1.25444492],  
[ 1.07292472],  
[ 1.03564825],  
[ 1.07940758],  
[ 1.17340911],  
[ 1.25120349],  
[ 1.45055157],  
[ 1.43758584],  
[ 1.08751116],  
[ 1.11506333],  
[ 1.32089428],  
[ 1.57210527],  
[ 1.9578357 ],  
[ 2.36625616],  
[ 2.26901319],  
[ 2.16042521],  
[ 1.87517918],  
[ 1.23013418],  
[ 1.17665055],  
[ 1.13775336],  
[ 0.96109531],  
[ 1.19934057],  
[ 1.2123063 ],  
[ 1.08913188],  
[ 1.1717884 ],  
[ 1.11506333],  
[ 1.18637484],  
[ 1.04861397],  
[ 0.56564058],  
[ 0.62398636],  
[ 0.25932524],  
[ 0.24960094],  
[ 0.10859864],  
[-0.47810058],  
[-0.59155071],  
[-0.47161772],  
[-0.74227731],  
[-0.7617259 ],  
[-0.80872667],  
[-0.92703894],  
[-0.88003818],  
[-0.75524304],  
[-0.60937859],
```

```

[-0.48782488],
[-0.48134202],
[-0.18475097],
[-0.11506018],
[-0.48134202],
[-0.53482565],
[-0.66934508],
[-1.05021337],
[-1.25442359],
[-1.12800774],
[-1.366253  ],
[-1.0437305  ],
[-0.62720647],
[-0.57858498],
[-0.58992999],
[-0.43109982],
[-0.02105865],
[ 0.17180657],
[ 0.05835644],
[ 0.01945926],
[-0.14099164],
[-0.12964662],
[ 0.00325209],
[-0.07778371],
[-0.1069566  ],
[ 0.05187358],
[-0.12802591],
[ 0.01783854],
[ 0.18639301],
[ 0.05187358]])

```

```

In [9]: print(X.shape)
        print (y.shape)
        y = y.reshape([y.shape[0],1])
        print (y.shape)

```

```

(148, 4, 2)
(148, 1)
(148, 1)

```

```

In [10]: X_train = X[:103,:,:]
        X_test = X[103,:,:]
        y_train = y[:103,:]
        y_test = y[103,:]
        print('X_train shape : ', X_train.shape)
        print('X_test shape : ', X_test.shape)
        print('y_train shape : ', y_train.shape)
        print('y_test shape : ', y_test.shape)

```

```

X_train shape : (103, 4, 2)
X_test shape : (45, 4, 2)
y_train shape : (103, 1)
y_test shape : (45, 1)

```

```
In [11]: # Initializing the Neural Network based on LSTM
model = Sequential()

# Adding 1st LSTM layer
model.add(LSTM(units=64, return_sequences=True, input_shape=(nlags, len(input_data))))

# Adding 2nd LSTM layer
model.add(LSTM(units=10, return_sequences=False))

# Adding Dropout
model.add(Dropout(0.2))

# Output layer
if(close_price_as_predictive_column):
    model.add(Dense(units=1, activation='linear'))
else:
    model.add(Dense(units=1, activation='tanh'))

# Compiling the Neural Network
if(close_price_as_predictive_column):
    model.compile(optimizer = 'adam', loss='mean_squared_error')
else:
    model.compile(loss='binary_crossentropy', optimizer='nadam', metrics=['accuracy'])
model.summary()
```

2023-05-20 03:48:44.799670: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 4, 64)	17152
lstm_1 (LSTM)	(None, 10)	3000
dropout (Dropout)	(None, 10)	0
dense (Dense)	(None, 1)	11

=====  
Total params: 20,163

Trainable params: 20,163

Non-trainable params: 0  
=====

```
In [12]: es = EarlyStopping(monitor='val_loss', min_delta=1e-10, patience=10, verbose
rlr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=10, verbose
mcp = ModelCheckpoint(filepath='weights.h5', monitor='val_loss', verbose=1,

tb = TensorBoard('logs')

history = model.fit(X_train, y_train, shuffle=True, epochs=50, callbacks=[es
```

Epoch 1/50

1/6 [====>.....] - ETA: 12s - loss: -0.0903 - accuracy: 0.0000e+00

Epoch 1: val\_loss improved from inf to -2.20356, saving model to weights.h5  
6/6 [=====] - 3s 126ms/step - loss: 3.1775 - accuracy: 0.0000e+00 - val\_loss: -2.2036 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 2/50

1/6 [====>.....] - ETA: 0s - loss: 3.2869 - accuracy: 0.0000e+00

Epoch 2: val\_loss did not improve from -2.20356  
6/6 [=====] - 0s 11ms/step - loss: 2.6256 - accuracy: 0.0000e+00 - val\_loss: -2.2036 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 3/50

1/6 [====>.....] - ETA: 0s - loss: -1.1279 - accuracy: 0.0000e+00

Epoch 3: val\_loss did not improve from -2.20356  
6/6 [=====] - 0s 12ms/step - loss: 1.7905 - accuracy: 0.0000e+00 - val\_loss: -2.2036 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 4/50

1/6 [====>.....] - ETA: 0s - loss: 0.3924 - accuracy: 0.0000e+00

Epoch 4: val\_loss improved from -2.20356 to -2.91778, saving model to weights.h5  
6/6 [=====] - 0s 13ms/step - loss: 1.8872 - accuracy: 0.0000e+00 - val\_loss: -2.9178 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 5/50

1/6 [====>.....] - ETA: 0s - loss: 3.0530 - accuracy: 0.0000e+00

Epoch 5: val\_loss did not improve from -2.91778  
6/6 [=====] - 0s 9ms/step - loss: 1.3304 - accuracy: 0.0000e+00 - val\_loss: -0.7277 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 6/50

1/6 [====>.....] - ETA: 0s - loss: 1.8208 - accuracy: 0.0000e+00

Epoch 6: val\_loss did not improve from -2.91778  
6/6 [=====] - 0s 13ms/step - loss: 0.9452 - accuracy: 0.0000e+00 - val\_loss: -2.3751 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 7/50

1/6 [====>.....] - ETA: 0s - loss: 1.8940 - accuracy: 0.0000e+00

Epoch 7: val\_loss did not improve from -2.91778  
6/6 [=====] - 0s 9ms/step - loss: 1.0320 - accuracy: 0.0000e+00 - val\_loss: -0.2386 - val\_accuracy: 0.0000e+00 - lr: 0.0010

Epoch 8/50

1/6 [====>.....] - ETA: 0s - loss: 1.4647 - accuracy: 0.

```
0000e+00
Epoch 8: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.7071 - accuracy
: 0.0000e+00 - val_loss: -0.1626 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 9/50
1/6 [====>.....] - ETA: 0s - loss: 1.1860 - accuracy: 0.
0000e+00
Epoch 9: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.6742 - accuracy
: 0.0000e+00 - val_loss: -0.1236 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 10/50
1/6 [====>.....] - ETA: 0s - loss: 0.9995 - accuracy: 0.
0000e+00
Epoch 10: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.6183 - accuracy
: 0.0000e+00 - val_loss: -0.1030 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 11/50
1/6 [====>.....] - ETA: 0s - loss: 0.7319 - accuracy: 0.
0000e+00
Epoch 11: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.5618 - accuracy
: 0.0000e+00 - val_loss: -0.0924 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 12/50
1/6 [====>.....] - ETA: 0s - loss: 0.7586 - accuracy: 0.
0000e+00
Epoch 12: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.6752 - accuracy
: 0.0000e+00 - val_loss: -0.0709 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 13/50
1/6 [====>.....] - ETA: 0s - loss: 0.2679 - accuracy: 0.
0000e+00
Epoch 13: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.6621 - accuracy
: 0.0000e+00 - val_loss: -0.0640 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 14/50
1/6 [====>.....] - ETA: 0s - loss: 0.7015 - accuracy: 0.
0000e+00
Epoch 14: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

Epoch 14: val_loss did not improve from -2.91778
6/6 [=====] - 0s 9ms/step - loss: 0.6302 - accuracy
: 0.0000e+00 - val_loss: -0.0620 - val_accuracy: 0.0000e+00 - lr: 0.0010
Epoch 14: early stopping
```

```
In [13]: X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 2))
X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 2))
model.evaluate(X_test, y_test)
predictions = model.predict(X_test)
predictions_train = model.predict(X_train)

if(close_price_as_predictive_column):
    y_pred = sc_predict.inverse_transform(predictions)
    y_pred_train = sc_predict.inverse_transform(predictions_train)
else:
    predictions[predictions >= 0] = 1
    predictions[predictions < 0] = -1
    predictions = predictions.astype('int32')

print(mean_squared_error(y_test,predictions))

2/2 [=====] - 0s 2ms/step - loss: 0.2761 - accuracy
: 0.0000e+00
2/2 [=====] - 0s 3ms/step
4/4 [=====] - 0s 2ms/step
2.2222222222222223
```

```
In [14]: from sklearn.metrics import classification_report
if(not close_price_as_predictive_column):
    print(classification_report(y_test, predictions))
else:
    plt.rcParams["figure.figsize"] = [10,5]
    plt.rcParams["figure.autolayout"] = True

    plt.plot(df.loc[:, 'Date'], df.loc[:, 'Close'], color='b', label='Actual S
    plt.plot(df.loc[:102, 'Date'], y_pred_train, color='orange', label='Train
    plt.plot(df.loc[103+nlags:, 'Date'], y_pred, color='r', label='Predicted

    plt.xticks(np.arange(0, len(df), 20))

    plt.legend(shadow=True)
    plt.title(company_name + ' (LO,Close price)',fontsize=12)
    plt.xlabel('Timeline', fontsize=10)
    plt.ylabel('Stock Price Value', fontsize=10)
    plt.show()
```

	precision	recall	f1-score	support
-1	0.33	0.04	0.07	24
1	0.45	0.90	0.60	21
accuracy			0.44	45
macro avg	0.39	0.47	0.34	45
weighted avg	0.39	0.44	0.32	45

In [ ]:

